

ITNPBD7 Assignment 2017

Movie Review Data

Your task in completing this assignment is to analyse some movie review data. The data are contained in a file called `ratedReviews.txt`, which you can download from the module assignments page (where you also found this document). The file contains the text of movie reviews, each followed by a score from 1 to 9 indicating a rating given to the movie. Each review and score pair are on a single line of the file and the scores are separated from the reviews by a tab character. Consequently, you can examine the file in Excel if you want.

Your task is to write the MapReduce code, in Java, needed to process the movie reviews in a way that discovers the most common word used in reviews with each possible rating. The result will be 9 words: one for each score, which represent the most common word used in reviews with each rating score. You should make sure common words such as *this*, *the*, *a* etc. are not counted.

You will submit a written report, detailing your design and the results you found. You will also be asked to submit a Java file containing your code.

Step 1, HDFS – 20 Marks

Before you write any code, you will need to copy the data onto your own space in HDFS. In your report, give details of how HDFS stores data such as this (assume the file is much bigger than it really is for the purpose of your description). This section should be around half a page long, plus a diagram. Describe what HDFS is for, the architecture it uses, and the roles of different nodes in the cluster. Document the `hdfs` commands you used to create a directory for the data and place it there.

Step 2, Design – 20 Marks

Now consider the MapReduce design you will implement. You know there are only nine different scores associated with the movie reviews and a larger (but unknown) number of different words used in those reviews. Consider and compare two different choices you could make to implement the given task. What keys and values will the mapper emit? Consider how much data will be moved across the network in each of your two designs. Also consider how many different reducers will be used in each case. Finally, choose one of the two designs to implement and justify your choice.

Step 3, Implement – 60 Marks

There is a list of words that should not be counted in the reviews – they are given in the file `Excl.txt`, which you can download from the assignments web page. You should now implement your design in Java using the Hadoop API that we have been using in class. Your code should find the most commonly used word (excluding those in the exclusion list) for movies with each rating from 1 to 9. You should allow the exclusion list to be supplied as a cache file; do not hard code the list into your Java. Make sure you implement a mapper, combiner and reducer.

What to submit

Please write up your work in a report, print it out and post it into the box marked ITNPBD7 outside 4B80. Additionally, please email your .java file to kms@cs.stir.ac.uk. The report should contain the content described above and a printout of your Java code. Make sure the Java code is very well commented – comments are worth 20 of the 60 marks associated with this part. Make sure your report also contains the results you got when you ran the code – that is the most common word for each rating category. The deadline for submission is Friday 31st of March at 4pm.